

# Data Access 2.0?

...please welcome...

# Spring Data!

Oliver Gierke



# Oliver Gierke

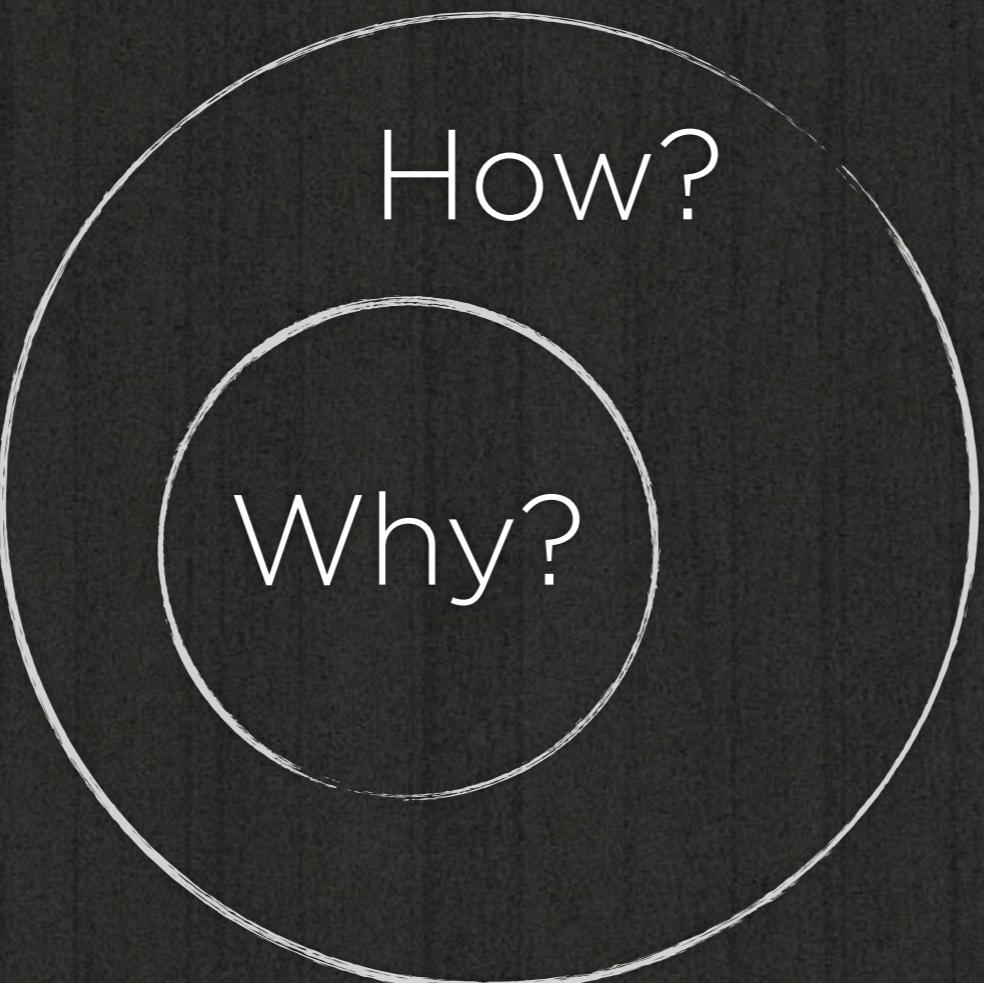


Spring Data  
Core/JPA/MongoDB  
JPA 2.1 Expert Group

- ✉ ogierke@gopivotal.com
- 🌐 www.olivergierke.de
- 🐦 olivergierke

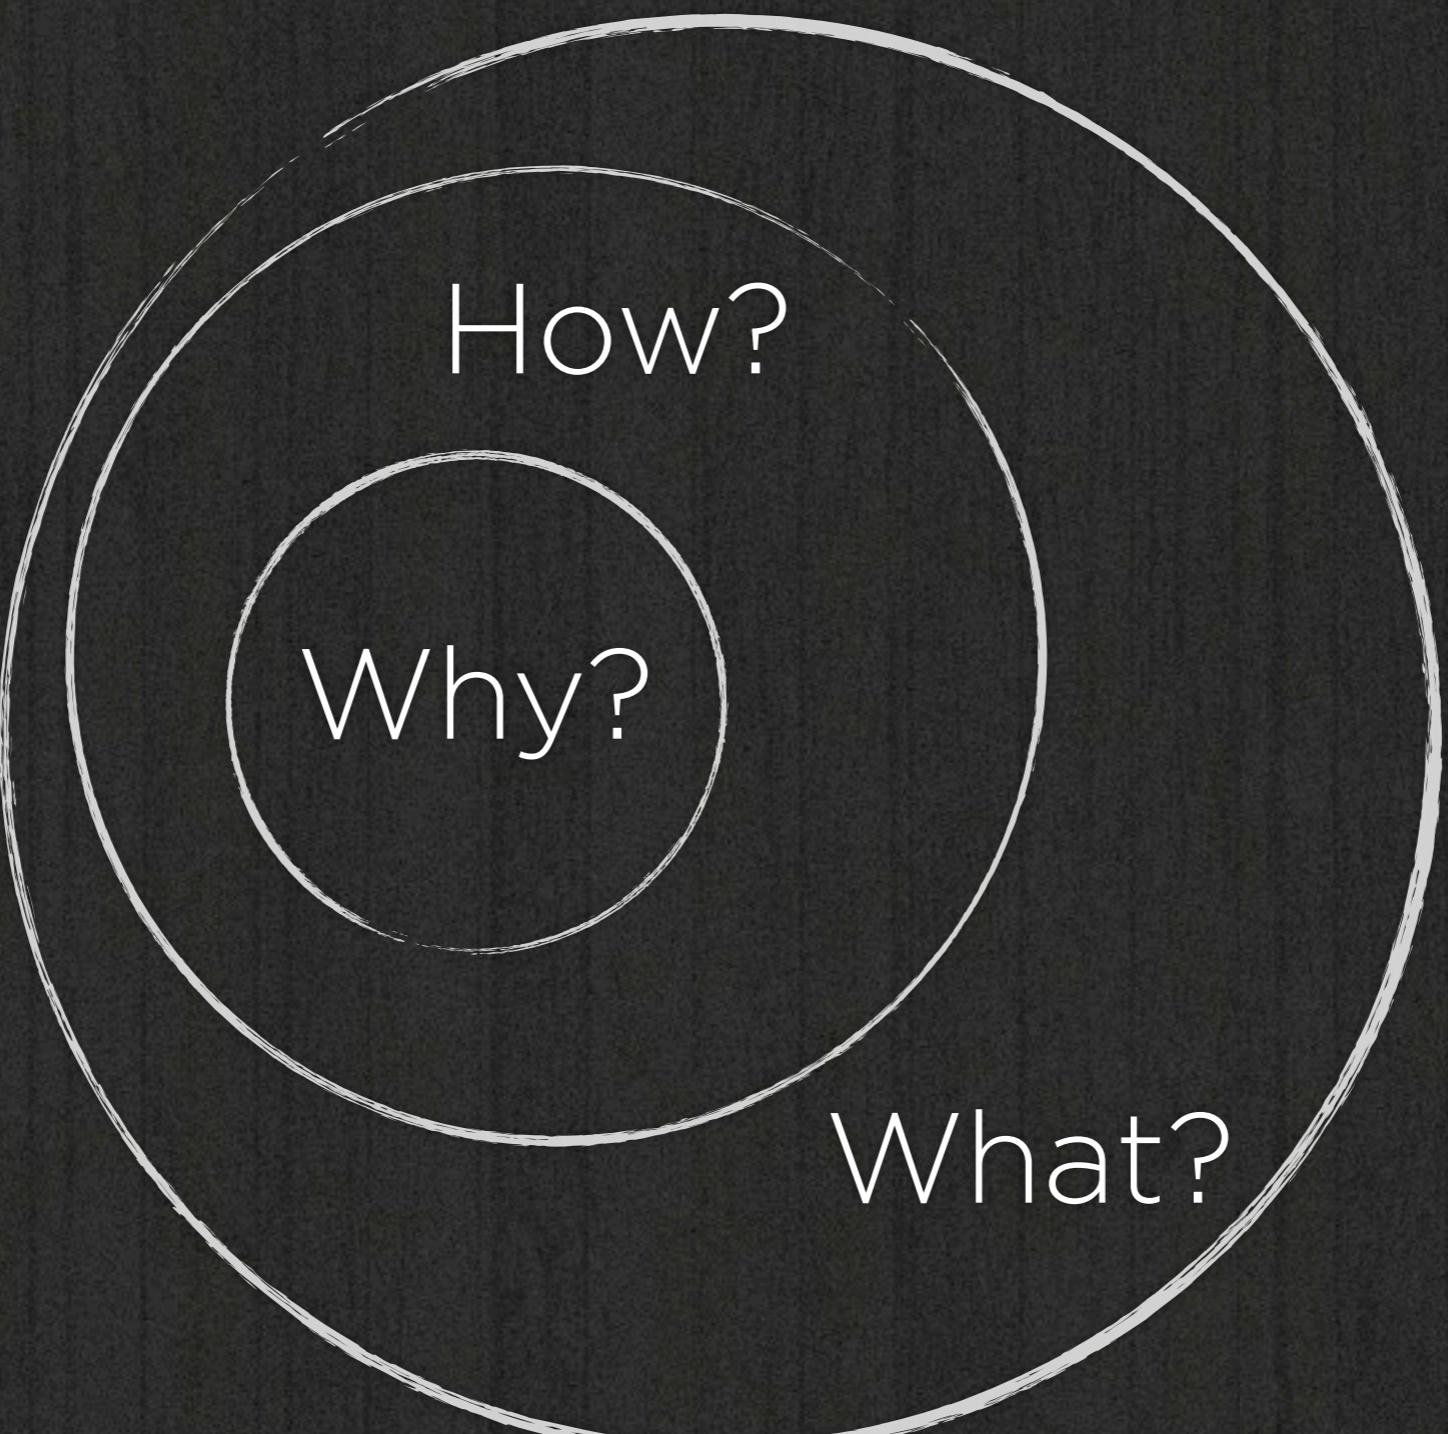
What to expect?

Why?



How?

Why?



How?

Why?

What?

# A Developer's View



What to expect?  
NOT!



14  
COLD BEER  
FROZEN PIZZAS  
FROZEN ROLLERS  
FROZEN VEGETABLES  
WATER

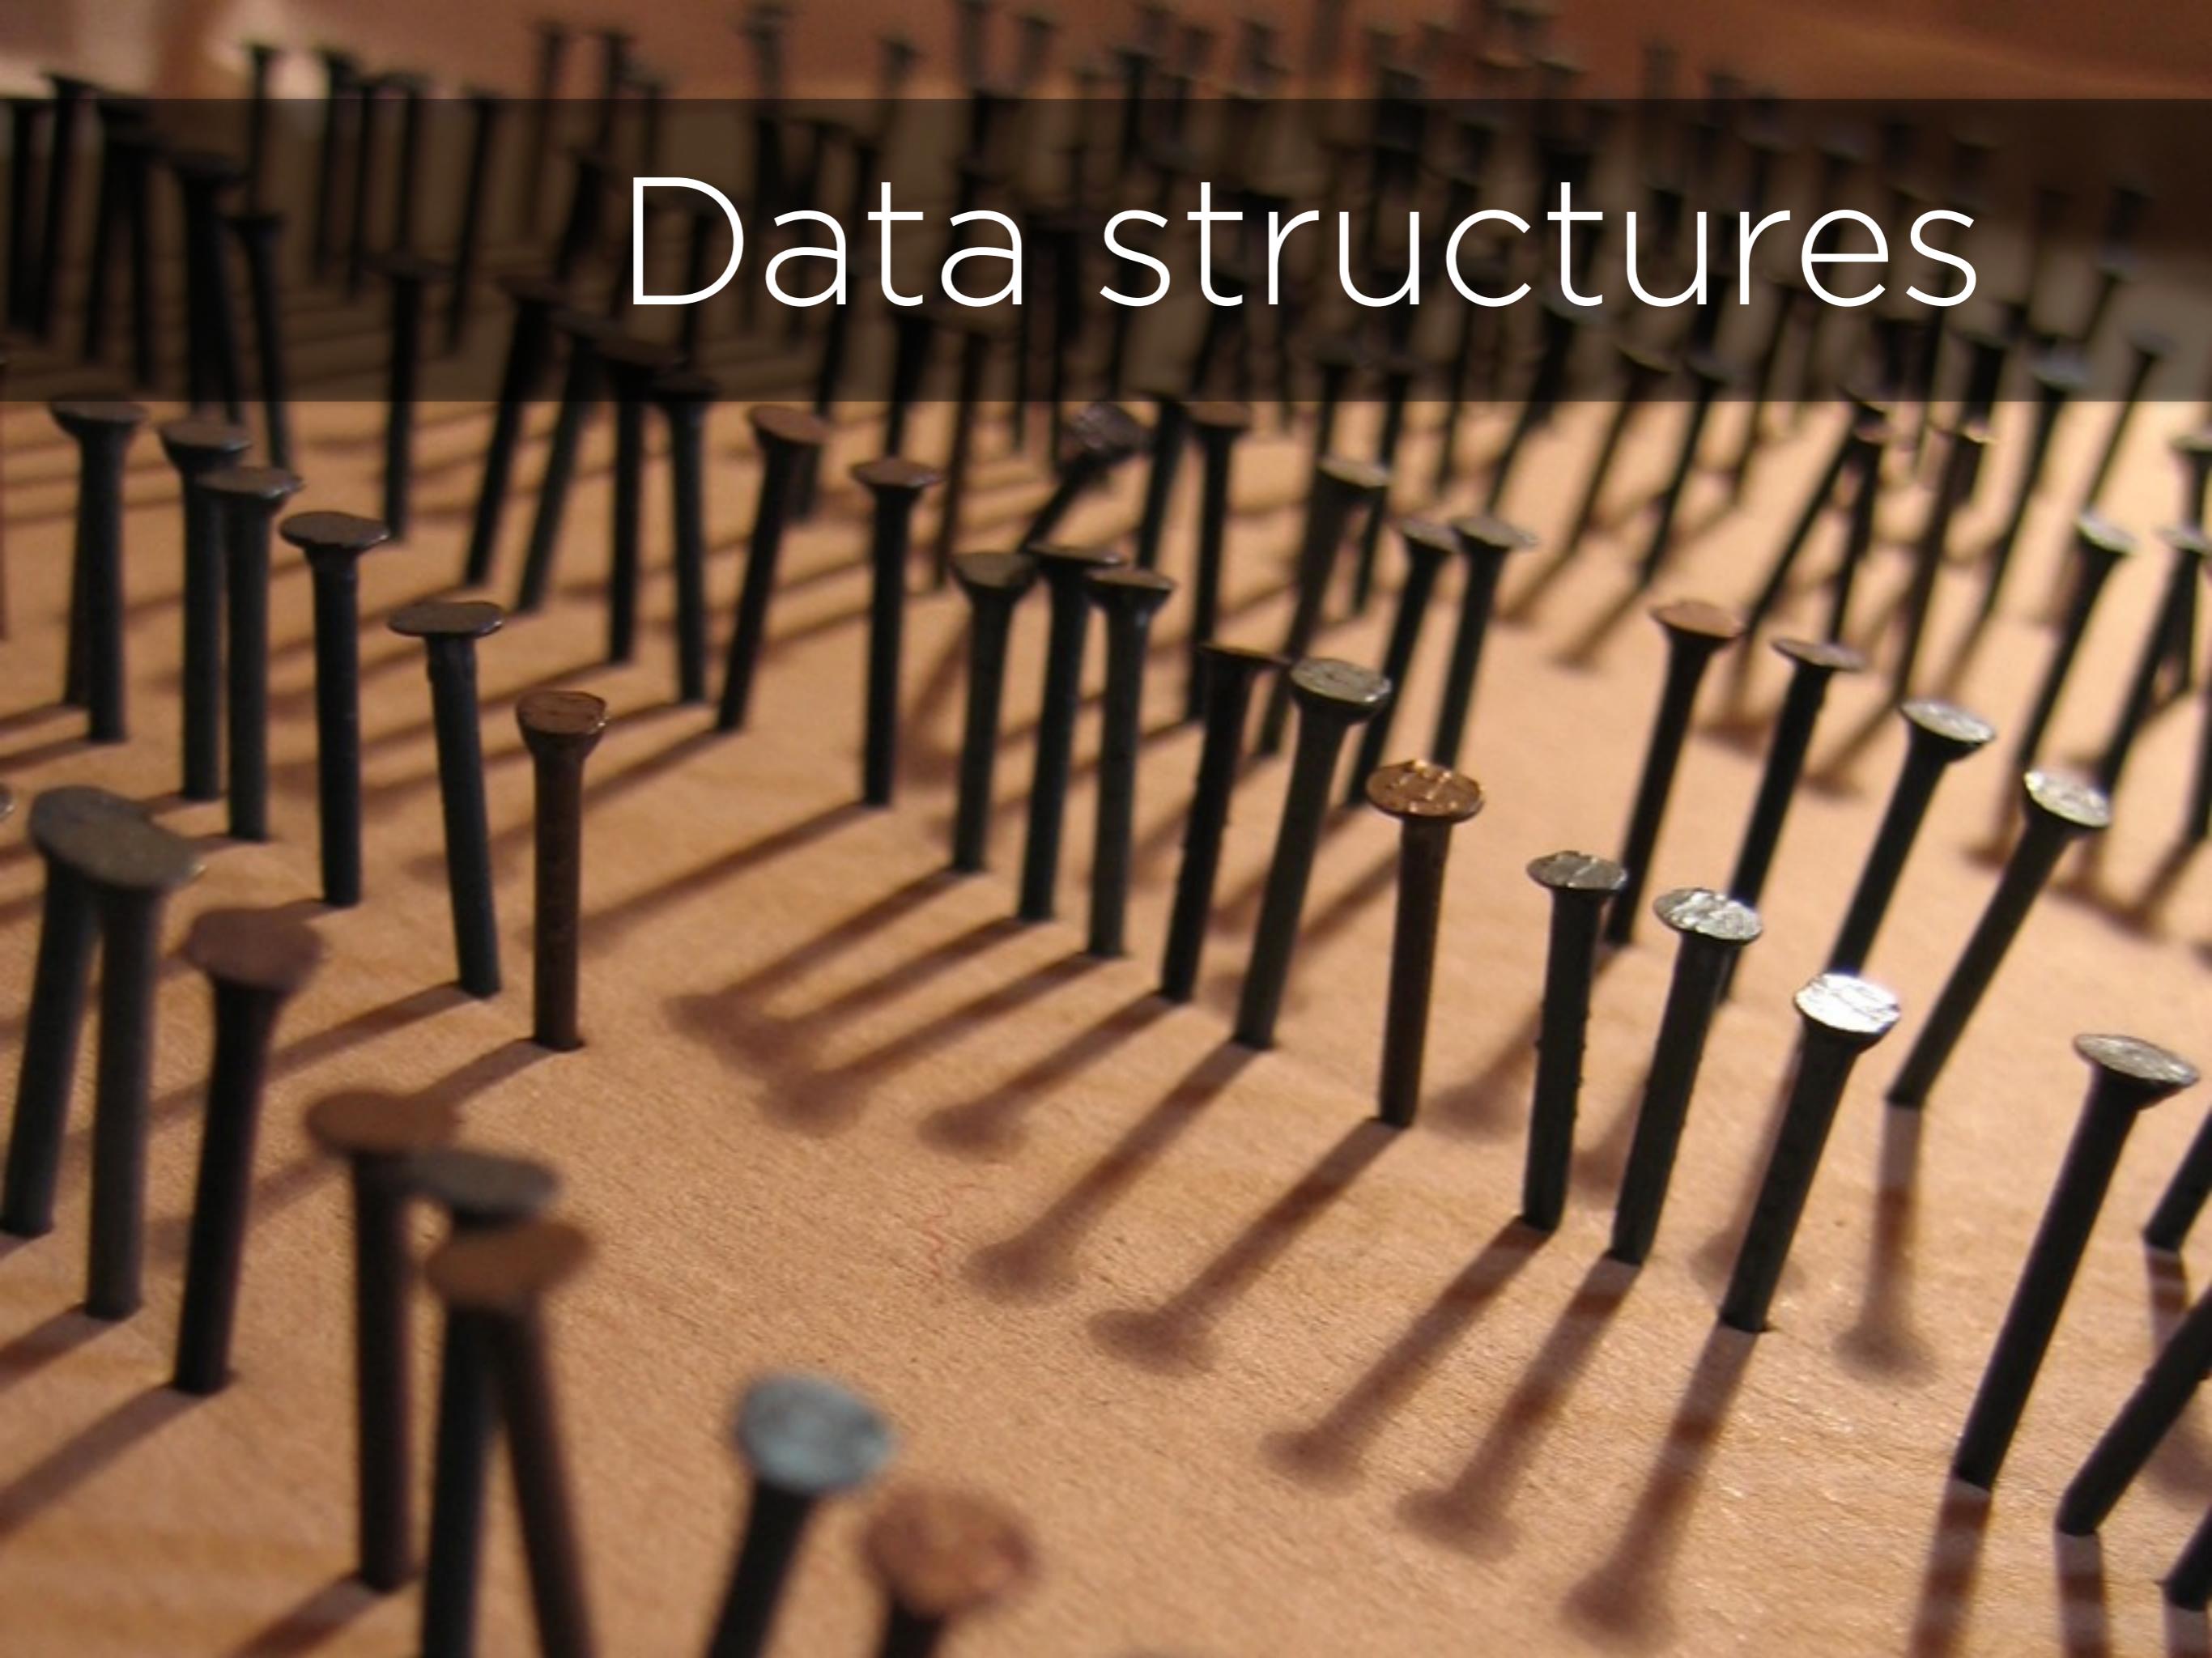
# Retrospect



# Relational databases

AIRCRAFT CATEGORY AND CLASS									
CAR 3005			ROUTE OF FLIGHT		TOTAL DURATION OF FLIGHT	AIRPLANE SINGLE-ENGINE LAND	AIRPLANE SINGLE-ENGINE SEA	AIRPLANE MULTI-ENGINE LAND	ROTORCRAFT HELICOPTER
DATE	AIRCRAFT MAKE AND MODEL	AIRCRAFT IDENT	FROM	TO					
3-22-05	PA144180	186WW	DUT	DUT	2.4			24	Seritzinger Lars 2729867 CFI Exp. 9-30-2005
04-01-05	PA-44-180	186WW	DVT-SEE-DVT		2.7				CHRIS MARRINGTON PA-44-180 2745805 CFI Exp.
04-06-05	PA-44-180	186WW	DVT-BLH-YUM-DVT		2.7				INGTON 2745805 CFI EXP
4-26-05	C-172-N	54715	DVT-SEL-DVT						
5-5-05	PA-44-180	215SE	DUT	DUT					
5-12-05	PA-44-180	3020M	DUT	DUT	1.6				
5-20-05	PA44180	3020M	DVT-PAN-DUT		1.8				
5-26-05	PA44180	215SE	DUT	DVI		7			
6-1-05	PA-44-180	215SE	DUT	DUT		6			
6-3-05	PA-44-180	215SE	DVT-ABQ-DUT		5.2				
6-8-05	PA-44-180	215SE	DUT-IMA-DUT		2.5				
6-13-05	PA-44-180	215SE	DVT-FFZ-DVT						
6-13-05	PA44180	146PA	DUT	DUT					
6-14-05	PA44180	215SE	DUT	10A-DUT	1	29	9		
TOTALS THIS PAGE									
AMT. FORWARDED					194	9	10		

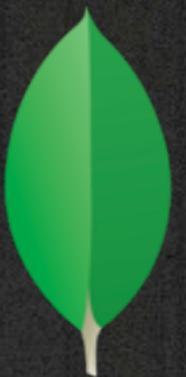
# Data structures



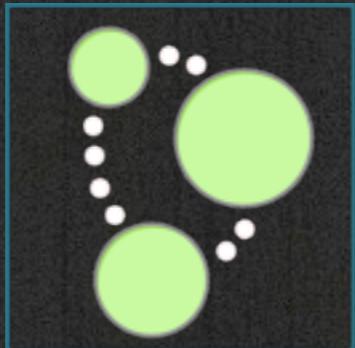


# Querying

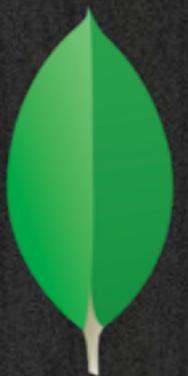
Hibari Voldemort  
Membase Riak Cassandra  
Redis SimpleDB (No)SQL MongoDB  
HBase OrientDB CouchDB  
Neo4J Couchbase



mongoDB



Neo4j  
the graph database



# mongoDB

Document database  
JSON documents  
JSON queries

# MongoDB Infrastructure API

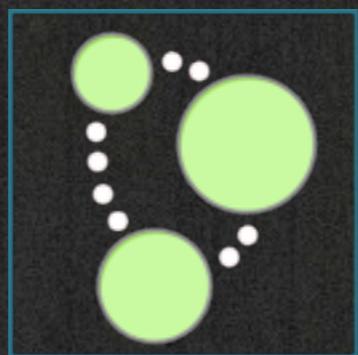
```
Mongo mongo = new Mongo(...);  
DB db = mongo.getDB("myDatabase");  
Collection collection = db.getCollection("myCollection");  
  
DBObject address = new BasicDBObject();  
address.put("city", "London");  
  
DBObject person = new BasicDBObject();  
person.put("firstname", "Dave");  
person.put("lastname", "Matthews");  
person.put("address", address);  
  
collection.save(person);
```



# MongoDB Query API

```
Mongo mongo = new Mongo(...);  
DB db = mongo.getDB("myDatabase");  
Collection collection = db.getCollection("myCollection");  
  
DBObject query = new BasicDBObject();  
query.put("address.city", "London");  
  
DBCursor cursor = collection.find(query);  
  
for (DBObject element : cursor) {  
    // Map data onto object  
}
```





**Neo4j**  
the graph database

Graph database  
Nodes / Relationships  
Traversals / Cypher / Gremlin

# Neo4J Infrastructure API

```
GraphDatabaseService database = new EmbeddedGraphDatabase(...);
Transaction tx = database.beginTx();

try {
    Node mrAnderson = database.createNode();
    mrAnderson.setProperty("name", "Thomas Anderson");
    Node morpheus = database.createNode();
    morpheus.setProperty("name", "Morpheus");

    Relationship friendship = mrAnderson.createRelationshipTo(
        morpheus, FriendTypes.KNOWS);

    tx.success();
} finally {
    tx.finish();
}
```



# Neo4J Query API

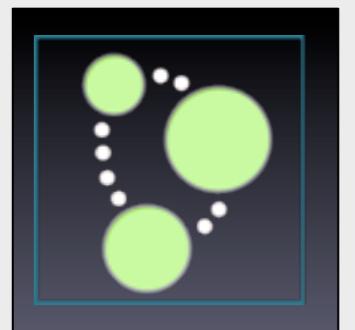
```
GraphDatabaseService database = new EmbeddedGraphDatabase(...);

CypherParser parser = new CypherParser();
Query query = parser.parse("start person = Person(id = *) match " +
    "person-[:colleagues]->colleague where colleague.firstname = {name}");

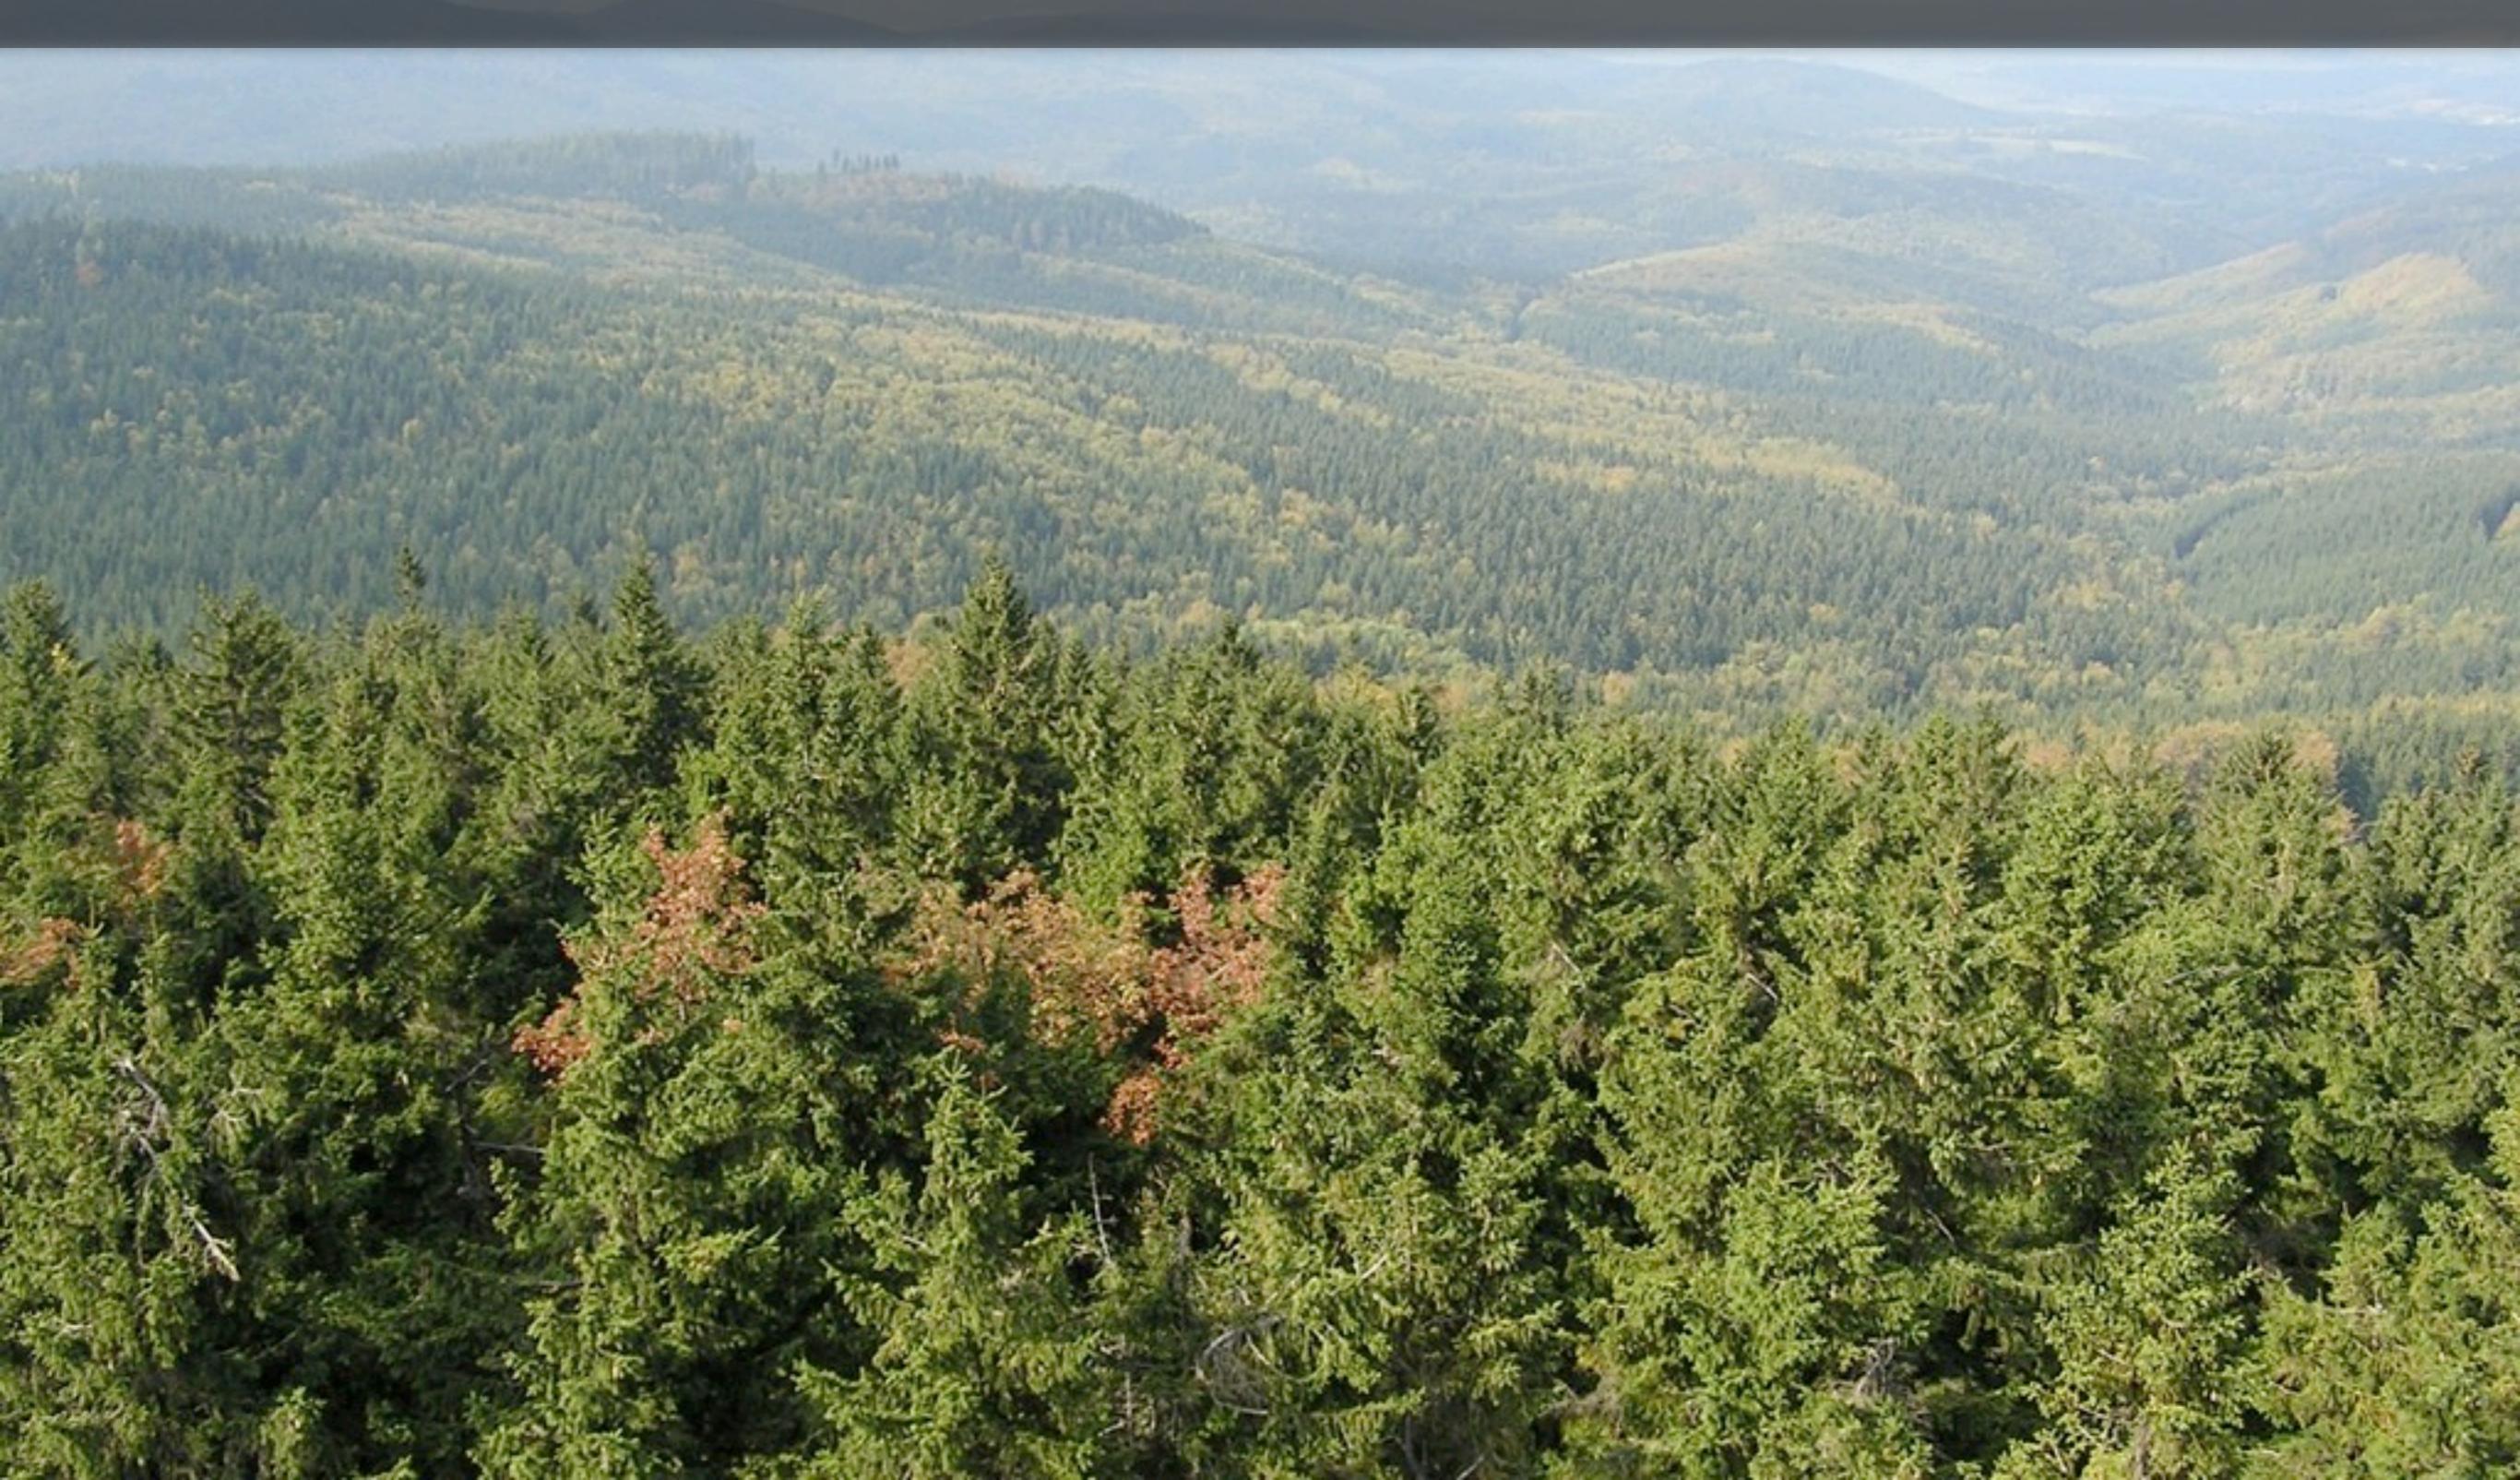
Map<String, Object> parameters = new HashMap<String, Object>();
parameters.put("name", "Dave");

ExecutionEngine engine = new ExecutionEngine(database);
ExecutionResult result = engine.execute(query, parameters);

for (EntrySet<String, Object> element : result) {
    // Map data onto object
}
```



# Forest for the woods?



JPA?

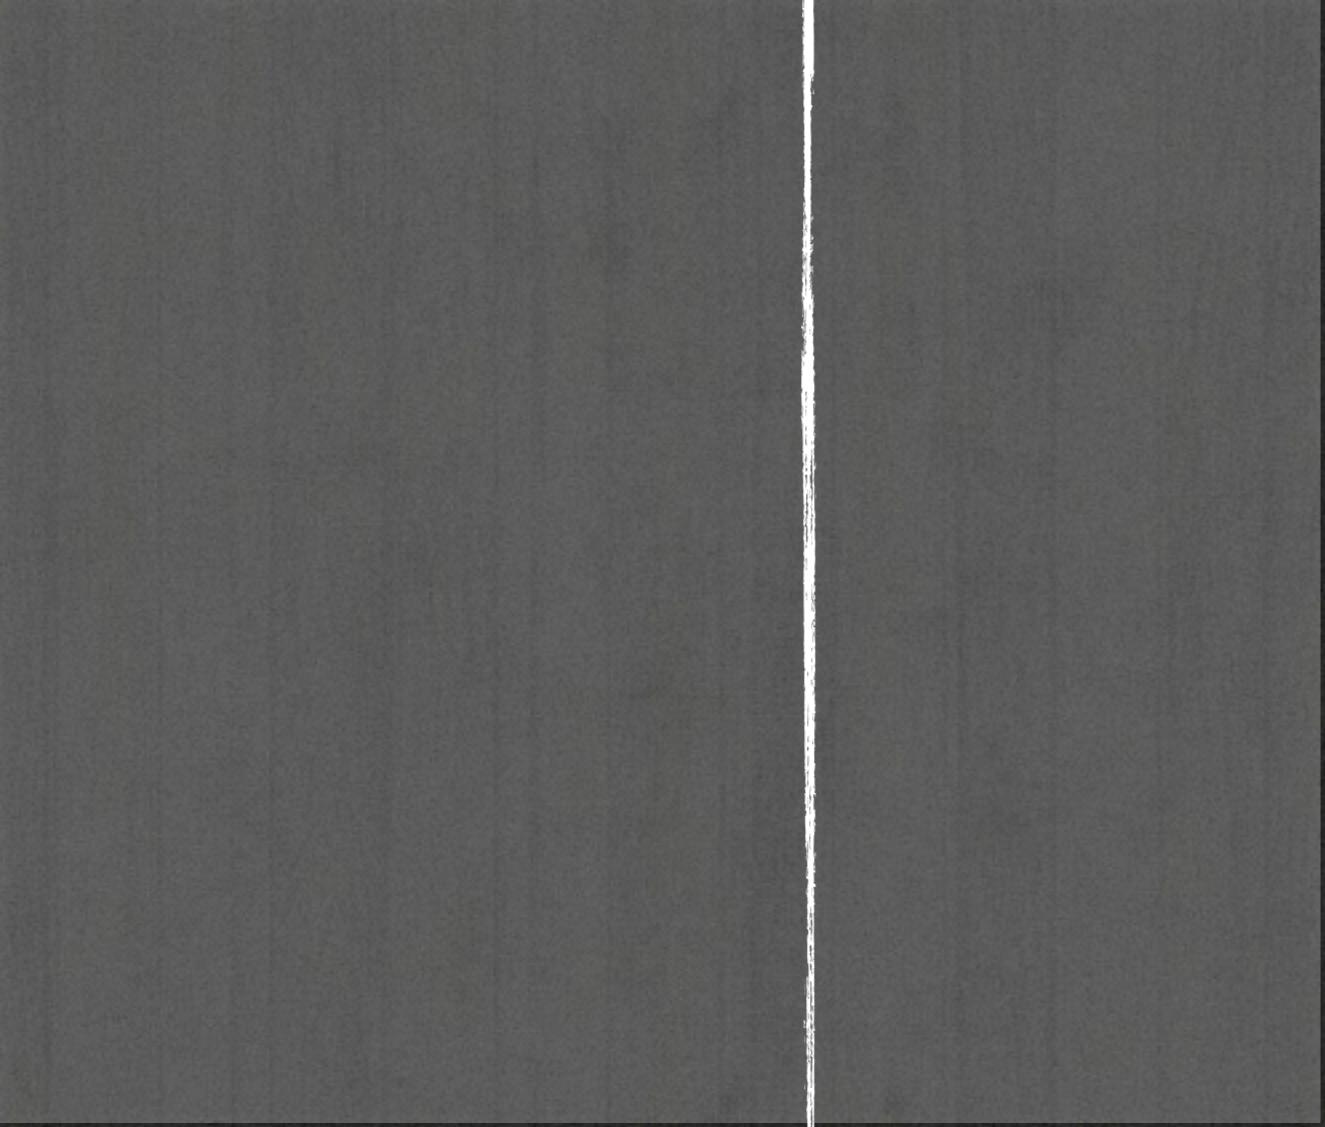
//

This document is the specification of the Java API for the management of persistence and object/relational mapping with Java EE and Java SE. The technical objective of this work is to provide an object/relational mapping facility for the Java application developer using a Java domain model to manage a relational database.

//

This document is the specification of the Java API for the management of persistence and **object/relational mapping** with Java EE and Java SE. The technical objective of this work is to provide an **object/relational mapping** facility for the Java application developer using a Java domain model to manage a relational database.

JPA



JPA



```
@Id  
@Entity
```

# JPA

@Table  
@Column

@Id  
@Entity

TX?

JPA

# JPA

@Table  
@Column

TX?

@Id  
@Entity

@Index  
@Field  
@RelationshipType

GeoSpatial  
Graph traversals

Store-specific  
functionality

# JPA

@Table  
@Column

TX?

@Id  
@Entity

@Index  
@Field  
@RelationshipType

GeoSpatial  
Graph traversals

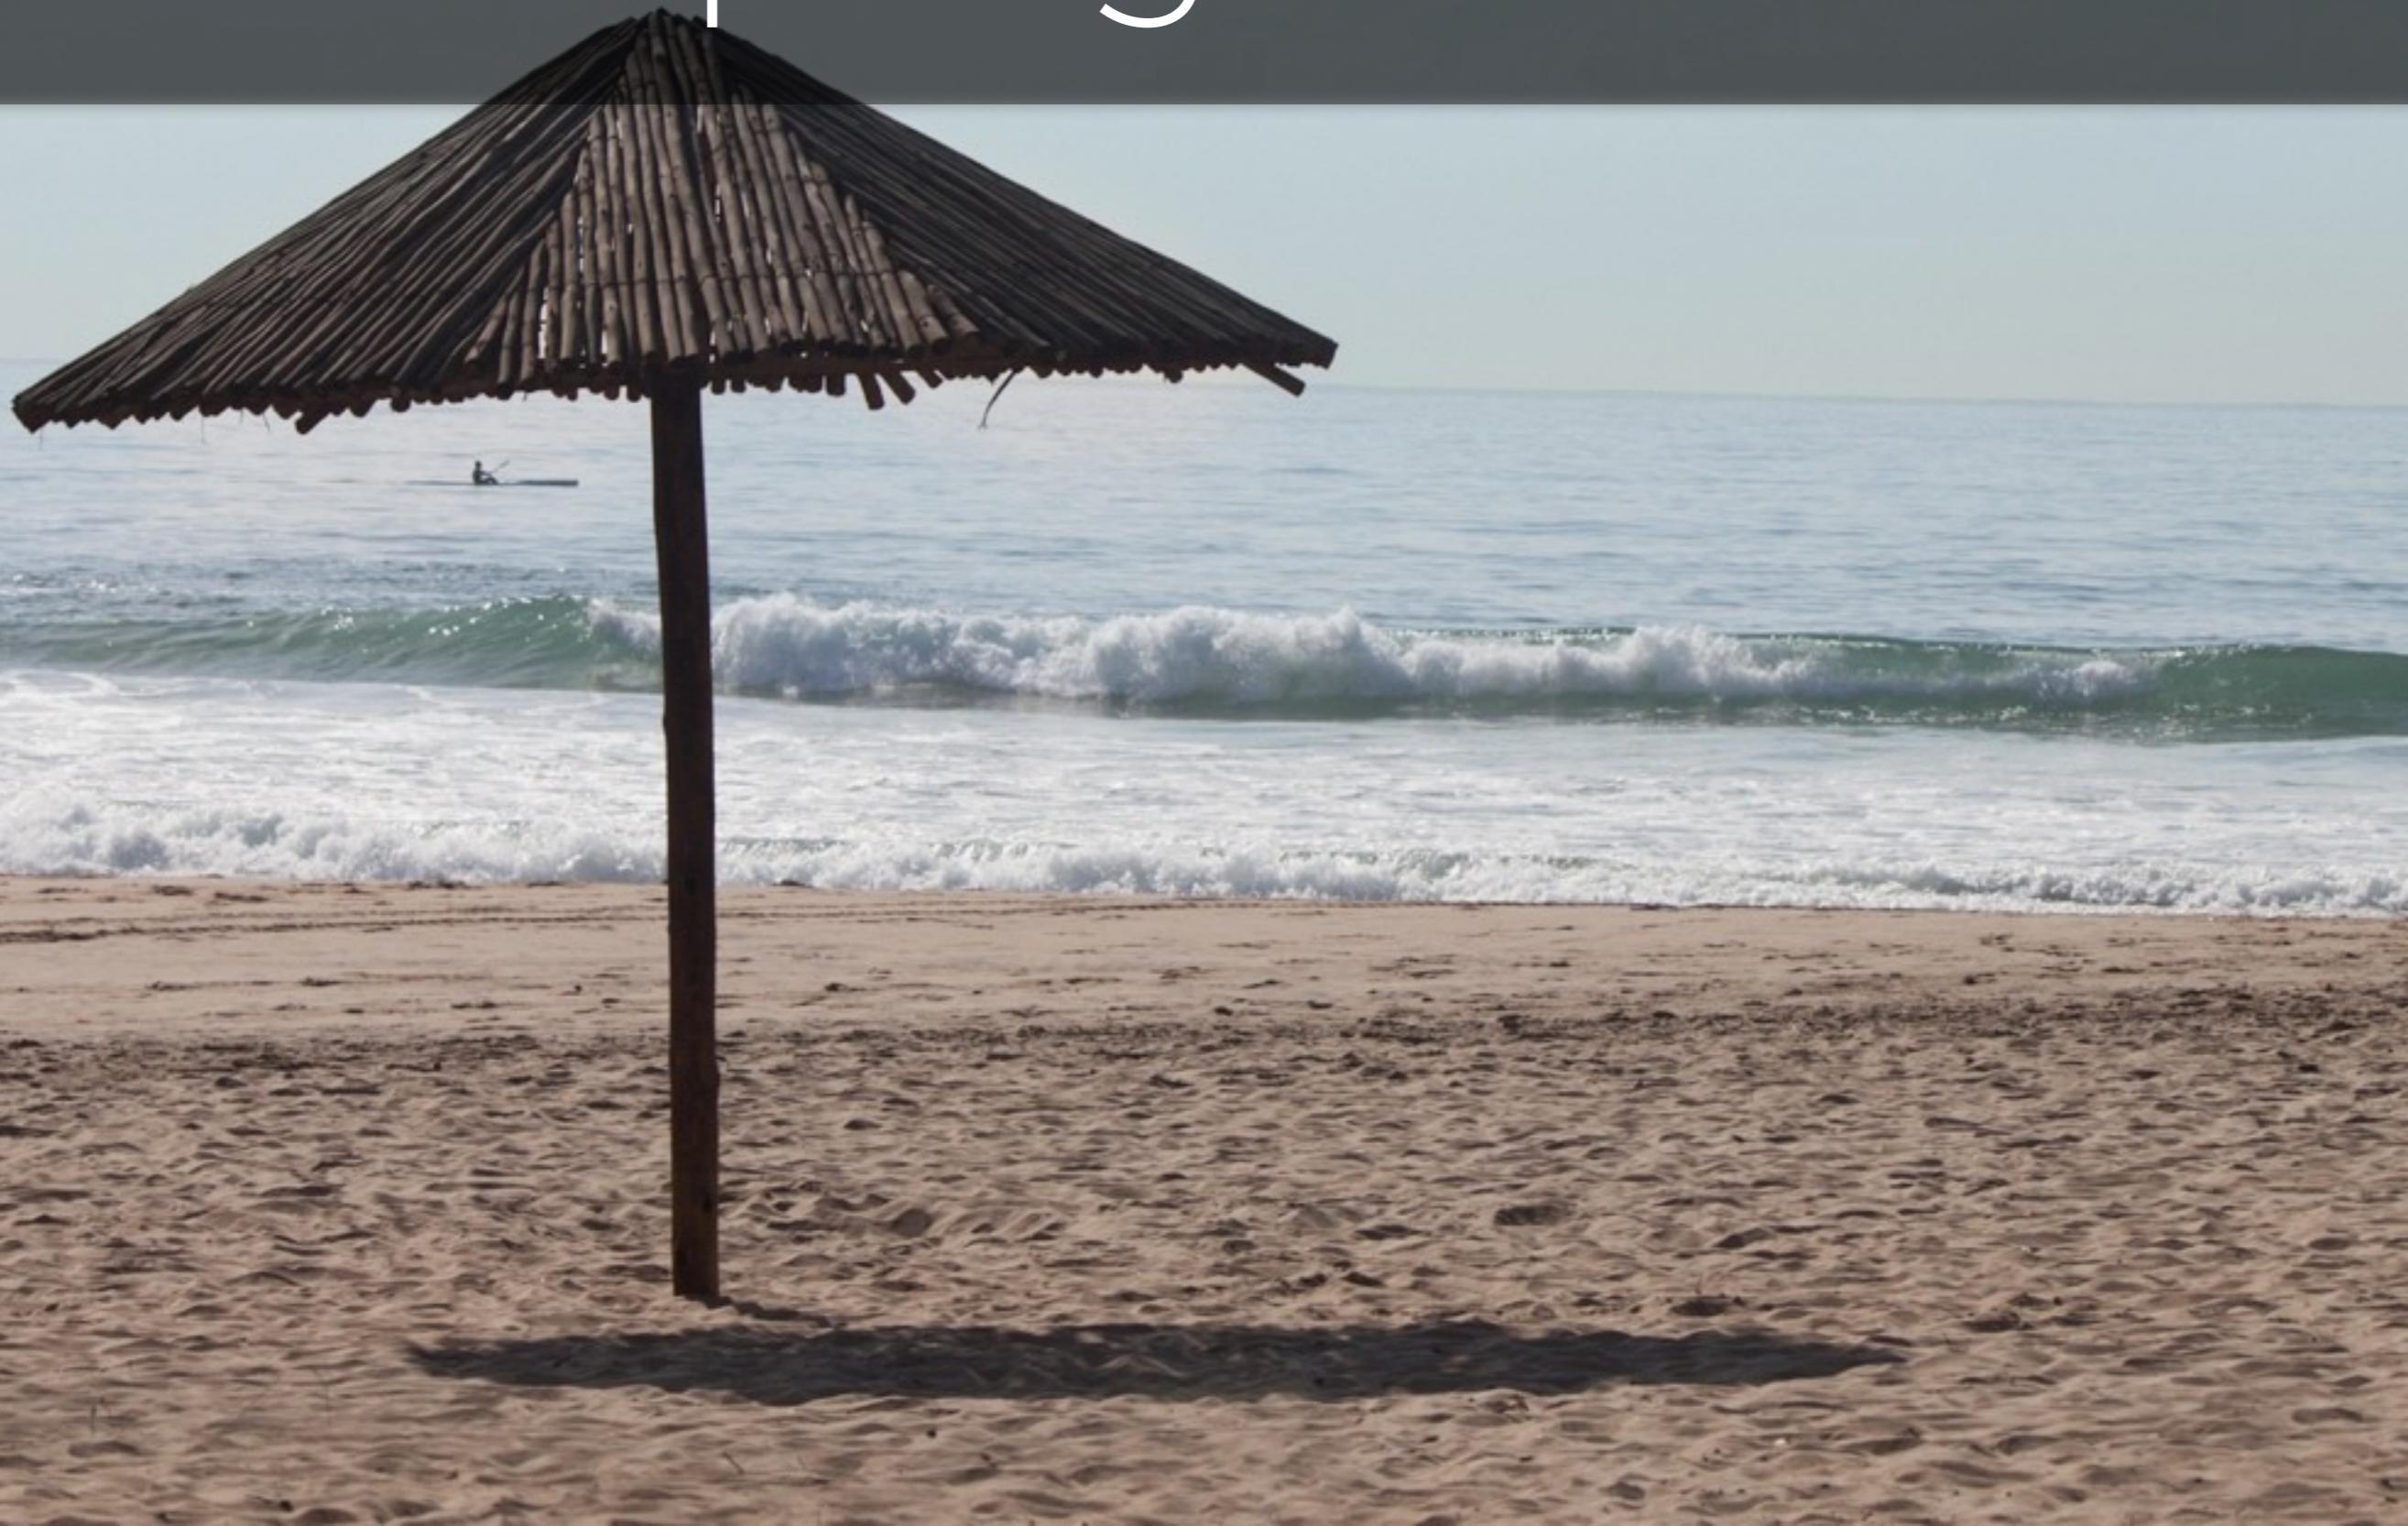
Store-specific  
functionality

JPA?

JPA?

There's some  
Spring for that!

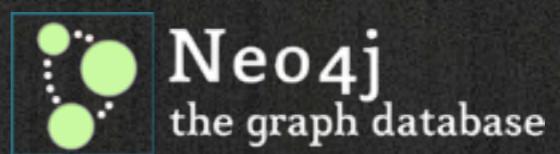
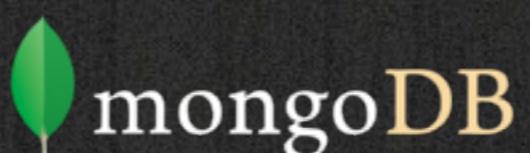
# Spring Data



!!!

... provide a familiar and consistent Spring-based programming model while retaining store-specific features and capabilities.

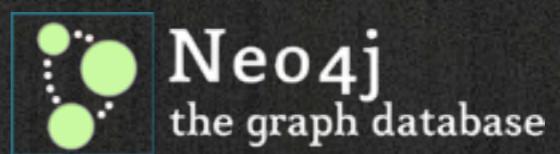
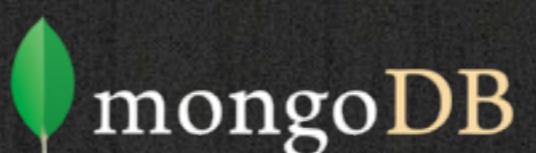
# Spring Data



elasticsearch.

JDBC / JPA

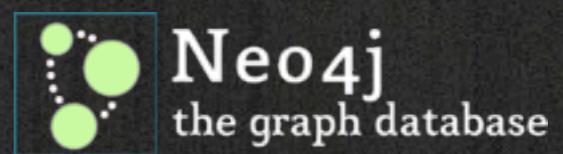
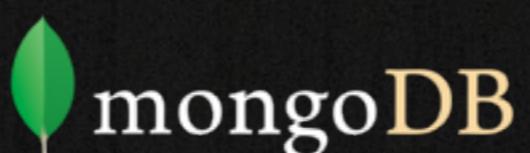
# Spring Data



elasticsearch.

JDBC / JPA

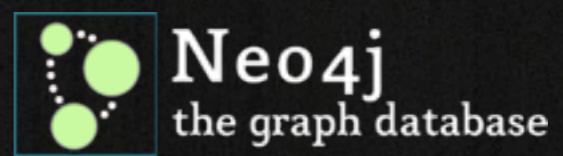
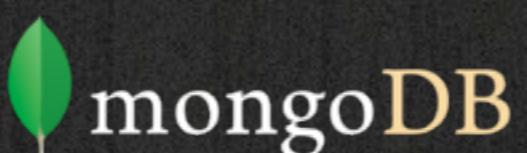
# Spring Data



elasticsearch.

JDBC / JPA

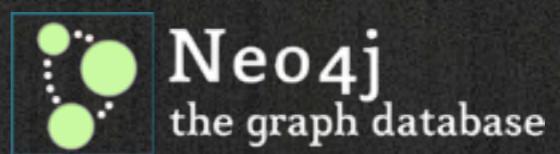
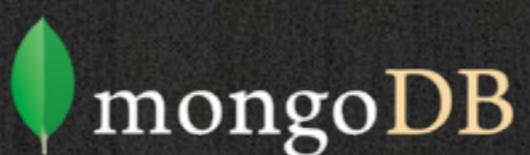
# Spring Data



elasticsearch.

JDBC / JPA

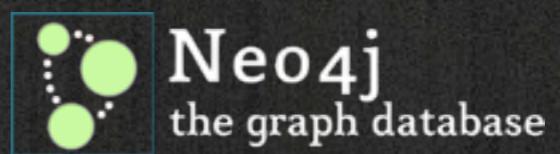
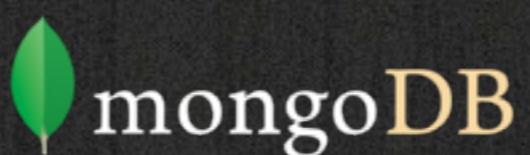
# Spring Data



elasticsearch.

JDBC / JPA

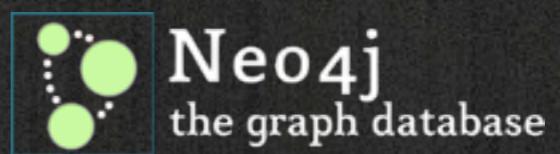
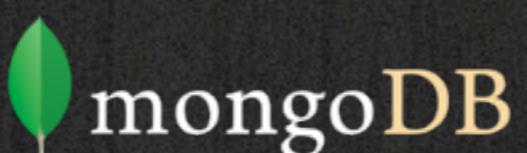
# Spring Data



elasticsearch.

JDBC / JPA

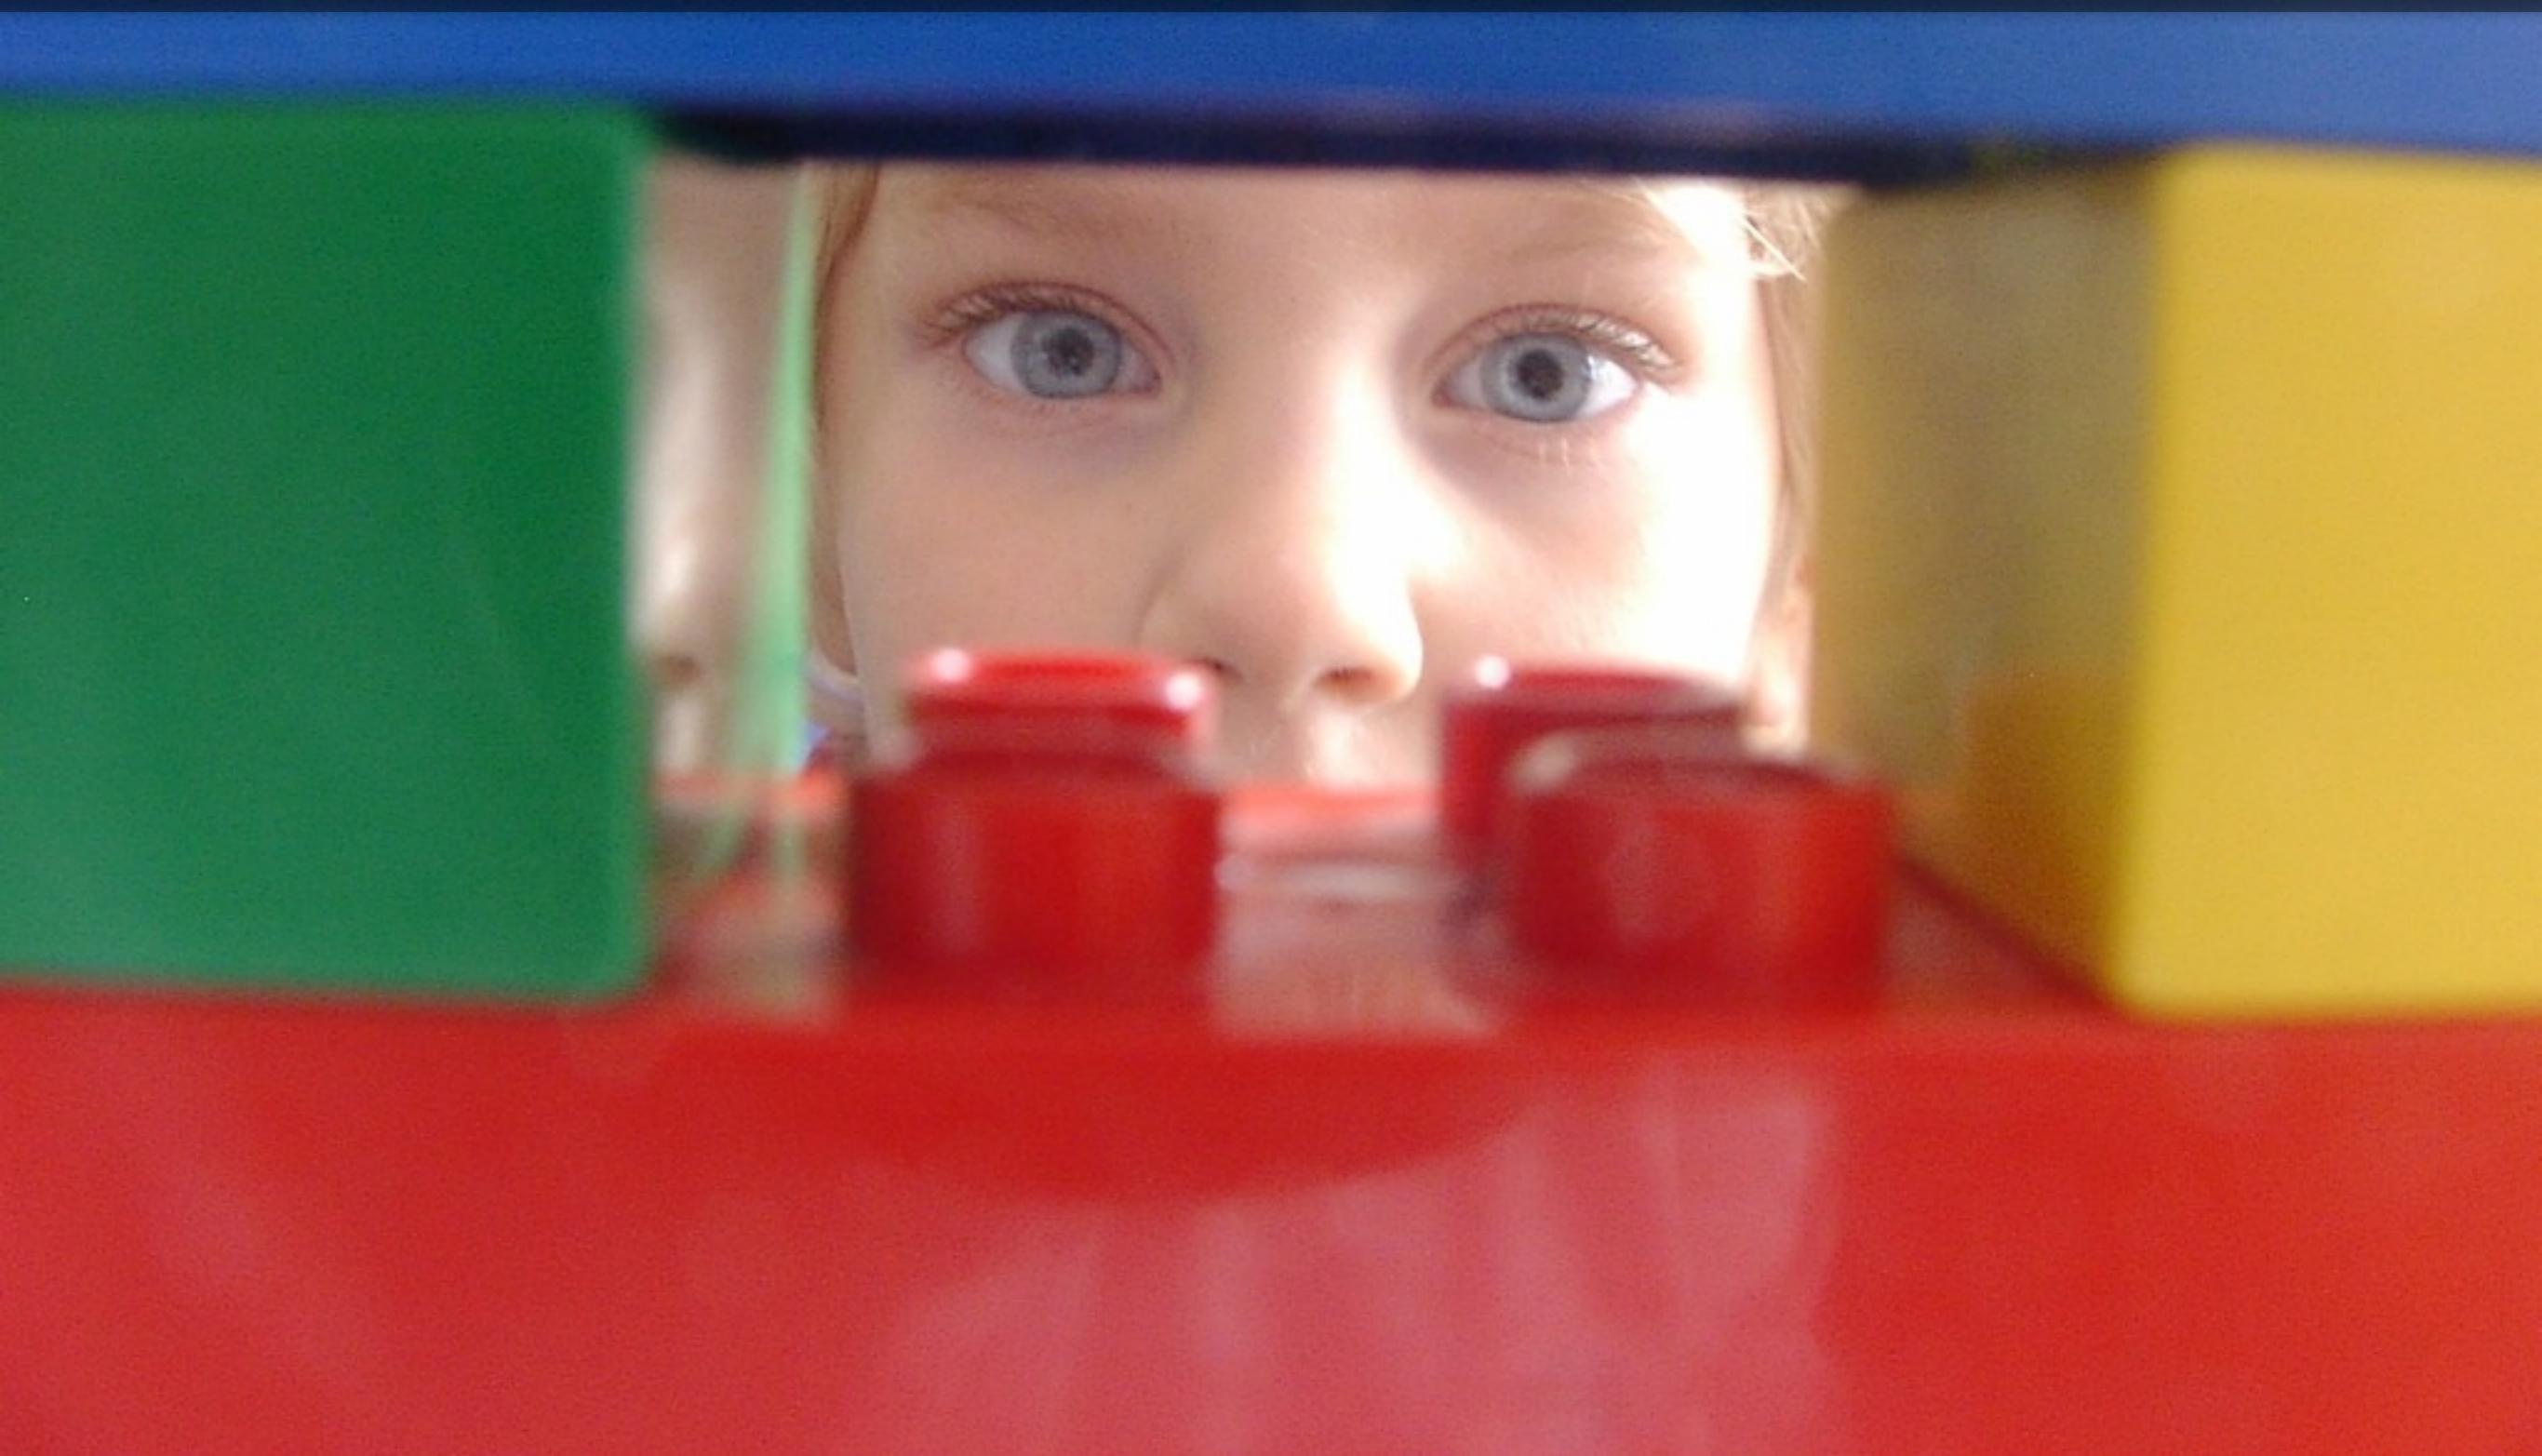
# Spring Data



elasticsearch.

JDBC / JPA

# Building blocks



# Spring



# Mapping



# JPA entity mapping

```
@Entity  
class Person {  
  
    @Id  
    @GeneratedValue(strategy=GenerationType.AUTO)  
    private BigInteger id;  
    private String firstname, lastname;  
  
    @Column(name="email")  
    private String emailAddress;  
  
    @OneToMany  
    private Set<Person> colleagues;  
}
```

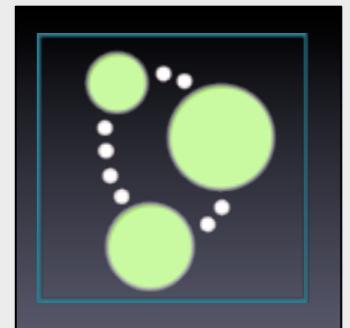
# Entity mapping - MongoDB

```
@Document  
class Person {  
  
    @Id private BigInteger id;  
    @Indexed private String firstname, lastname;  
    @Field("email") private String emailAddress;  
    @DBRef private Set<Person> colleagues;  
  
    public Person(String firstname) { ... }  
  
    @PersistenceConstructor  
    public Person(String firstname, String lastname) { ... }  
  
    ...  
}
```



# Entity mapping - Neo4J

```
@NodeEntity  
class Person {  
  
    @GraphId private long id;  
    @Indexed private String firstname, lastname;  
    @RelatedTo(direction = Direction.INCOMING)  
    private Set<Person> colleagues;  
  
    public Person(String firstname) { ... }  
  
    @PersistenceConstructor  
    public Person(String firstname, String lastname) { ... }  
  
    ...  
}
```



# Templates



# MongoOperations / -Template

```
public interface MongoOperations {  
  
    // Generic callback-accepting methods  
    <T> T execute(DbCallback<T> action);  
    <T> T execute(Class<?> entityClass, CollectionCallback<T> action);  
    <T> T execute(String collectionName, CollectionCallback<T> action);  
  
    // Higher level access methods  
    <T> List<T> find(Query query, Class<T> entityClass);  
    void save(Object objectToSave, String collectionName);  
    WriteResult updateFirst(Query query, Update update, Class<?>  
                           entityClass);  
  
    // Geo API  
    <T> GeoResults<T> geoNear(NearQuery near, Class<T> entityClass);  
}
```



# MongoTemplate usage

```
// Setup infrastructure
```

```
Mongo mongo = new Mongo();
MongoDbFactory factory = new SimpleMongoDbFactory(mongo, „foo“);
MongoTemplate template = new MongoTemplate(factory);
```

```
// Create and save entity
```

```
Person dave = new Person("Dave", "Matthews");
dave.setEmailAddress("dave@dmband.com");
template.save(person);
```

```
// Query entity
```

```
Query query = new Query(new Criteria("emailAddress")
                           .is("dave@dmband.com"));
assertThat(template.findOne(query, Person.class), is(dave));
```



# Repositories



# Repositories - JPA

```
<jpa:repositories base-package="com.acme.repositories" />

public interface PersonRepository extends Repository<Person, BigInteger>
{
    // Finder for a single entity
    Person findByEmailAddress(String emailAddress);

    // Finder for multiple entities
    List<Person> findByLastnameLike(String lastname);

    // Finder with pagination
    Page<Person> findByFirstnameLike(String firstname, Pageable page);
}
```

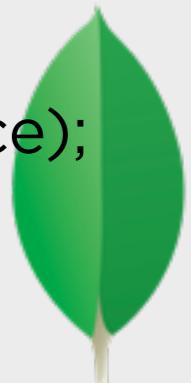
# Repositories - MongoDB

```
public interface PersonRepository extends Repository<Person, BigInteger>
{
    // Finder for a single entity
    Person findByEmailAddress(String emailAddress);

    // Finder for multiple entities
    List<Person> findByLastnameLike(String lastname);

    // Finder with pagination
    Page<Person> findByFirstnameLike(String firstname, Pageable page);

    // Geospatial queries
    List<Person> findByLocationNear(Point location, Distance distance);
    GeoResults<Person> findByLocationNear(Point location);
}
```



# Repositories - MongoDB

```
<mongo:repositories base-package="com.acme.repositories" />

@Component
public class MyClient {

    @Autowired
    private PersonRepository repository;

    public List<Person> doSomething() {

        Point point = new Point(43.7, 48.8);
        Distance distance = new Distance(200, Metrics.KILOMETERS);
        return repository.findByLocationNear(point, distance);
    }
}
```



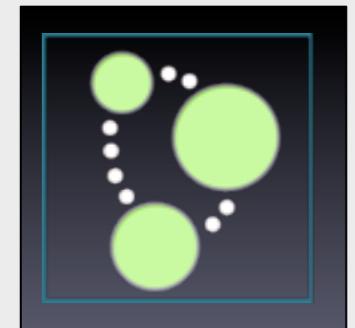
# Repositories - Neo4J

```
interface PersonRepository extends GraphRepository<Person, Long>
    // Finder for a single entity
    Person findByEmailAddress(String emailAddress);

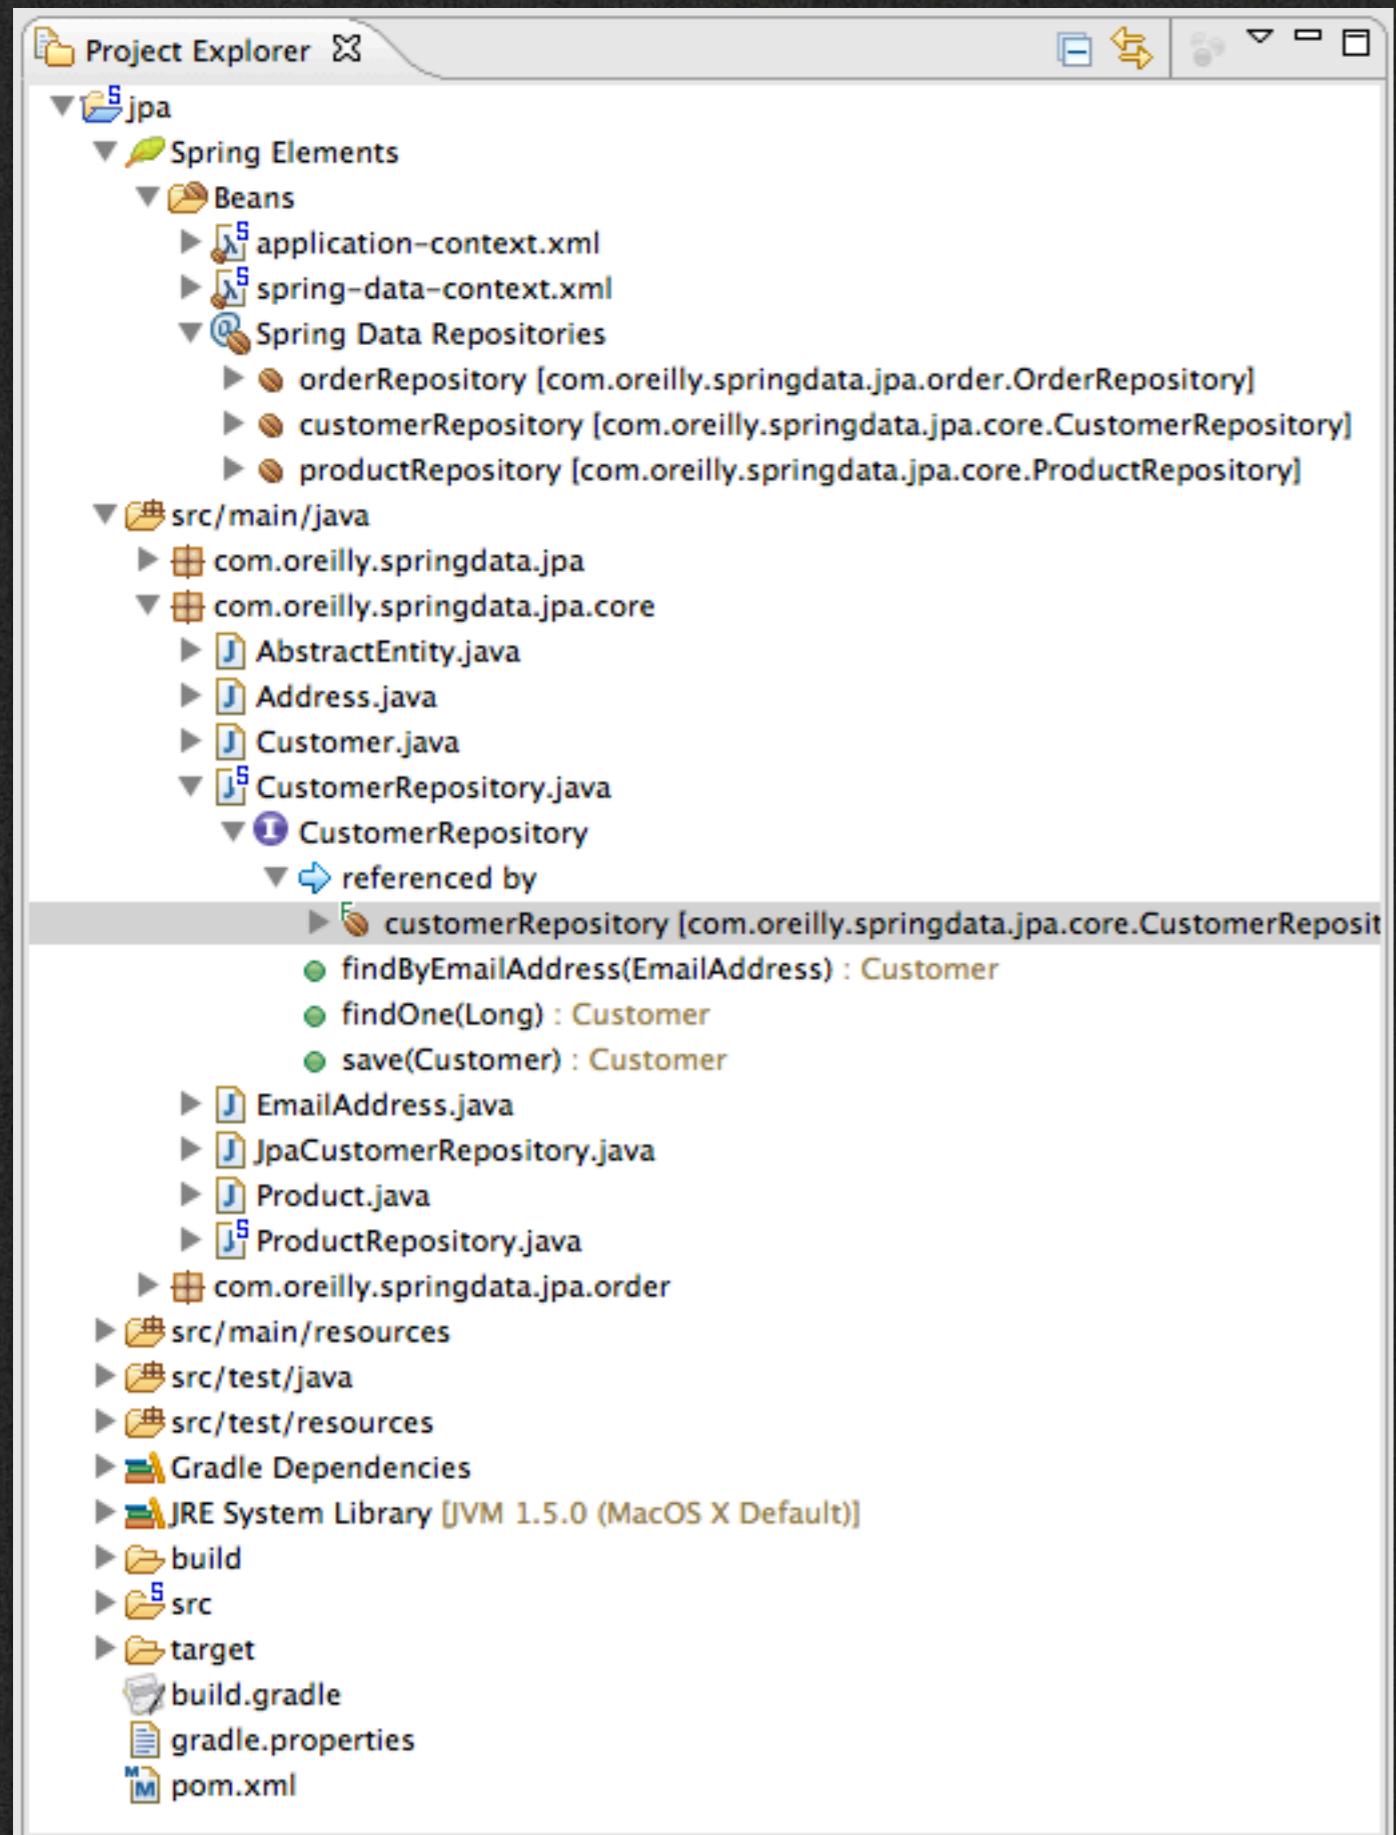
    // Finder for multiple entities
    List<Person> findByLastnameLike(String lastname);

    // Finder with pagination
    Page<Person> findByFirstnameLike(String firstname, Pageable page);

    @Query("start person = Person(id = *) " +
        "match person-[:colleagues]->colleague where " +
        "colleague.firstname = {name}")
    List<Person> getPersonsWithColleaguesName(
        @Param("name") Movie m);
}
```







Project Explorer

jpa

Spring Elements

Beans

application-context.xml

spring-data-context.xml

@ Spring Data Repositories

orderRepository [com.oreilly.springdata.jpa.core]

customerRepository [com.oreilly.springdata.jpa.core]

productRepository [com.oreilly.springdata.jpa.core]

src/main/java

com.oreilly.springdata.jpa

com.oreilly.springdata.jpa.core

AbstractEntity.java

Address.java

Customer.java

CustomerRepository.java

CustomerRepository

referenced by

customerRepository (com.oreilly.springdata.jpa.core.CustomerRepository)

findByEmailAddress(EmailAddress) : Customer

findOne(Long) : Customer

save(Customer) : Customer

EmailAddress.java

JpaCustomerRepository.java

Product.java

ProductRepository.java

com.oreilly.springdata.jpa.order

src/main/resources

src/test/java

src/test/resources

Gradle Dependencies

JRE System Library [JVM 1.5.0 (MacOS X Default)]

build

src

target

build.gradle

gradle.properties

pom.xml

ProductRepository.java

```
* Copyright 2012 the original author or authors.
package com.oreilly.springdata.jpa.core;

import org.springframework.data.domain.Page;
import org.springframework.data.domain.Pageable;
import org.springframework.data.repository.CrudRepository;

/**
 * Repository to access {@link Product} instances.
 *
 * @author Oliver Gierke
 */
public interface ProductRepository extends CrudRepository<Product, Long> {

    Page<Product> findByDescriptionContaining(String description, Pageable pageable);
}
```

Invalid derived query! No property description found for type com.oreilly.springdata.jpa.core.Product

The screenshot displays a Java development environment with two code editors open:

- ProductRepository.java** (Top Editor):

```
* Copyright 2012 the original author or authors.  
package com.oreilly.springdata.jpa.core;  
  
import org.springframework.data.domain.Page;  
import org.springframework.data.domain.Pageable;  
import org.springframework.data.repository.CrudRepository;  
  
orderRepository [com.oreilly.springdata.jpa.order.OrderRepository]  
customerRepository [com.oreilly.springdata.jpa.core.CustomerRepository]
```
- \*Order.java** (Bottom Editor):

```
/**  
 * An order.  
 *  
 * @author Oliver Gierke  
 */  
@Entity  
@Table(name = "Orders")  
public class Order extends AbstractEntity {  
  
    @ManyToOne  
    private Customer customer;  
    @ManyToOne  
    private Address billingAddress;  
    @ManyToOne  
    private Address shippingAddress;  
  
    @OneToMany(cascade = CascadeType.ALL, orphanRemoval = true)  
    @JoinColumn(name = "order_id")  
    private Set<LineItem> lineItems = new HashSet<LineItem>();  
  
    // ...  
}
```

The Project Explorer on the left shows the project structure:

- src
- build
- target
- build.gradle
- gradle.properties
- pom.xml
- Gradle Dependencies
- JRE System Library [JVM 1.5.0 (MacOS X Default)]

The screenshot shows an IDE interface with three code editors open:

- ProductRepository.java**:

```
* Copyright 2012 the original author or authors.  
package com.oreilly.springdata.jpa.core;  
  
import org.springframework.data.domain.Page;  
import org.springframework.data.domain.Pageable;  
import org.springframework.data.repository.CrudRepository;  
  
orderRepository [com.oreilly.springdata.jpa.core.OrderRepository]  
customerRepository [com.oreilly.springdata.jpa.core.CustomerRepository]
```
- \*Order.java**:

```
/**  
 * An order.  
 *  
 * @author Oliver Gierke  
 */  
@Entity  
@Table(name = "Orders")  
public class Order extends AbstractEntity  
  
    @ManyToOne  
    private Customer customer;  
    @ManyToOne  
    private Address billingAddress;  
    @ManyToOne  
    private Address shippingAddress;  
  
    @OneToMany(cascade = CascadeType.ALL, or  
    @JoinColumn(name = "order_id")  
    private Set<LineItem> lineItems = new He  
  
    // ...
```
- \*OrderRepository.java**:

```
/**  
 * Repository to access {@link Order}s.  
 *  
 * @author Oliver Gierke  
 */  
public interface OrderRepository extends PagingAndSortingRepository<Order, Long> {  
  
    List<Order> findByCustomer(Customer customer);  
}
```

A code completion dropdown is visible, listing properties of the Order class: **billingAddress**, **customer**, **lineItems**, **shippingAddress**, **shippingAddress**, **lineItems**, **customer**.

The Project Explorer on the left shows the project structure:

- src
- build
- target
- build.gradle
- gradle.properties
- pom.xml

Other visible files in the Project Explorer include application-context.xml, spring-data-context.xml, and various XML configuration files.

The screenshot shows an IDE interface with three code editors open:

- ProductRepository.java**:

```
* Copyright 2012 the original author or authors.  
package com.oreilly.springdata.jpa.core;  
  
import org.springframework.data.domain.Page;  
import org.springframework.data.domain.Pageable;  
import org.springframework.data.repository.CrudRepository;  
  
orderRepository [com.oreilly.springdata.jpa.order.OrderRepository]  
customerRepository [com.oreilly.springdata.jpa.core.CustomerRepository]
```
- \*Order.java**:

```
/**  
 * An order.  
 *  
 * @author Oliver Gierke  
 */  
@Entity  
@Table(name = "Orders")  
public class Order extends AbstractEntity  
  
    @ManyToOne  
    private Customer customer;  
    @ManyToOne  
    private Address billingAddress;  
    @ManyToOne  
    private Address shippingAddress;  
  
    @OneToMany(cascade = CascadeType.ALL, or  
    @JoinColumn(name = "order_id")  
    private Set<LineItem> lineItems = new He  
  
    // ...
```
- \*OrderRepository.java**:

```
/**  
 * Repository to access {@link Order}s.  
 *  
 * @author Oliver Gierke  
 */  
public interface OrderRepository extends PagingAndSortingRepository<Order, Long> {  
  
    List<Order> findByBillingAddress(Customer customer);  
}
```

The Project Explorer on the left shows the project structure with files like application-context.xml, spring-data-context.xml, and various beans defined.

A code completion dropdown is visible on the right side of the OrderRepository.java editor, listing various Spring Data annotations and methods.



Querydsl

# Querydsl

```
QPerson $ = QPerson.person;
BooleanExpression left = $.lastname.contains("eth");
BooleanExpression right = $.firstname.is("Carter");

public interface QueryDslPredicateExecutor<T> {
    T findOne(Predicate predicate);
    List<T> findAll(Predicate predicate);
}

public interface PersonRepository extends Repository<Person, BigInteger>,
    QueryDslPredicateExecutor { ... }

List<Person> result = repository.findAll(left.or(right));
assertThat(result.size(), is(2));
assertThat(result, hasItems(dave, carter));
```



Wrap up

# Mapping support

# Templates

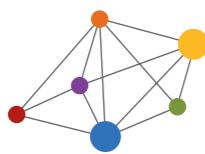
# Repositories

Store-specific  
functionality

QueryDSL

# CDI Integration

IDE support



## CONTENTS INCLUDE:

- › About the Spring Data Project
- › Configuration Support
- › Object Mapping
- › Template APIs
- › Repositories
- › Advanced Features... and more!

## Core Spring Data

*By: Oliver Gierke*

### ABOUT THE SPRING DATA PROJECT

The Spring Data project is part of the ecosystem surrounding the Spring Framework and constitutes an umbrella project for advanced data access related topics. It contains modules to support traditional relational data stores (based on plain JDBC or JPA), NoSQL ones (like MongoDB, Neo4j or Redis), and big data technologies like Apache Hadoop. The core mission of the project is to provide a familiar and consistent Spring-based programming model for various data access technologies while retaining store-specific features and capabilities.

### General Themes

#### Infrastructure Configuration Support

A core theme of all the Spring Data projects is support for configuring resources to access the underlying technology. This support is implemented using XML namespaces and support classes for Spring JavaConfig allowing you to easily set up access to a Mongo database, an embedded Neo4j instance, and the like. Also, integration with core Spring functionality like JMX is provided, meaning that some stores will expose statistics through their native API, which will be exposed to JMX via Spring Data.

#### Object Mapping Framework

Most of the NoSQL Java APIs do not provide support to map domain objects onto the stores' data model (e.g., documents in MongoDB, or nodes and relationships for Neo4j). So, when working with the native Java drivers, you would usually have to write a significant amount of code to map data onto the domain objects of your application when reading, and vice versa on writing. Thus, a core part of the Spring Data project is a mapping and conversion API that allows obtaining metadata about domain classes to be persisted and enables the conversion of arbitrary domain objects into store-specific data types.

#### Template APIs

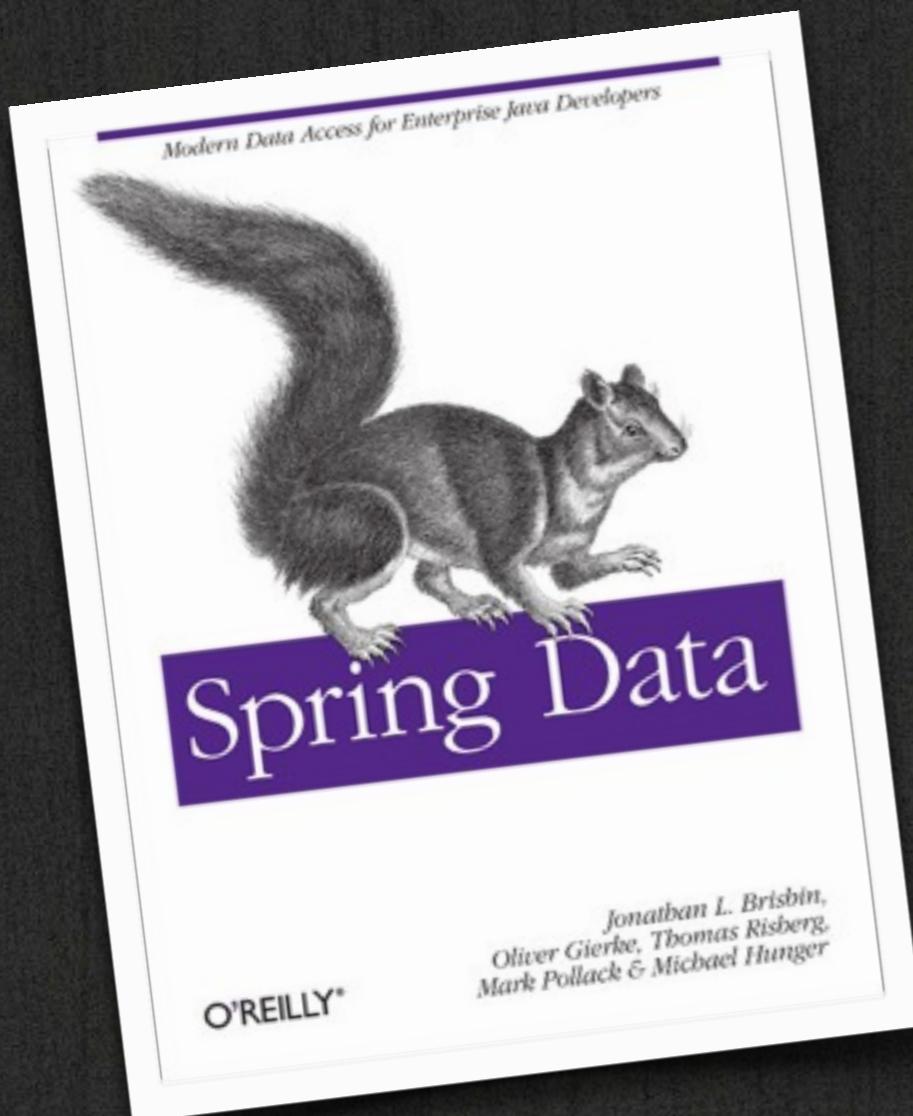
### JPA

XML element	Description
<jpa:repositories />	Enables Spring Data repositories support for repository interfaces underneath the package configured in the base-package attribute. JavaConfig equivalent is @EnableJpaRepositories.
<jpa:auditing />	Enables transparent auditing of JPA managed entities. Note that this requires the AuditingEntityListener applied to the entity (either globally through a declaration in orm.xml or through @EntityListener on the entity class).

### MongoDB

For Spring Data MongoDB XML namespace elements not mentioning a dedicated @Enable annotation alternative, you usually declare an @Bean-annotated method and use the plain Java APIs of the classes that would have otherwise been set up by the XML element. Alternatively, you can use the JavaConfig base class AbstractMongoConfiguration that Spring Data MongoDB ships for convenience.

XML element	Description
<mongo:db-factory />	One stop shop to set up a Mongo instance pointing to a particular database instance. For advanced-use cases define a <mongo:mongo /> externally and refer to it using a mongo-ref attribute.
<mongo:mongo />	Configures a Mongo instance. Supports basic attributes like host, port, write concern etc. Configure more advanced options through the nested <mongo:options /> element. In JavaConfig simply declare an @Bean



# Spring Data

Modern Data Access For Enterprise Java

JPA      JDBC      Hive  
NoSQL      Big Data      Pig  
MongoDB      Hadoop      HBase  
Redis      Gemfire      Splunk  
Querydsl      Neo4j      REST exporter  
Repositories

# Questions?



# Resources

- [springio.org/spring-data](http://spring.io/org/spring-data)
- [github.com/spring-projects/spring-data-jpa](https://github.com/spring-projects/spring-data-jpa)
- [github.com/spring-projects/spring-data-mongodb](https://github.com/spring-projects/spring-data-mongodb)
- [github.com/spring-projects/spring-data-neo4j](https://github.com/spring-projects/spring-data-neo4j)
- [se-radio.net/2010/07/episode-165-nosql-and-mongodb-with-dwight-merriman](http://se-radio.net/2010/07/episode-165-nosql-and-mongodb-with-dwight-merriman)
- [kkovacs.eu/cassandra-vs-mongodb-vs-couchdb-vs-redis](http://kkovacs.eu/cassandra-vs-mongodb-vs-couchdb-vs-redis)
- [github.com/spring-projects/spring-data-jpa-examples](https://github.com/spring-projects/spring-data-jpa-examples)